



CMA-tuned recurrent baselines converge to a common FLOPs-per-bit rate (N=4 model families, ~480M parameters each)

A. Loss vs FLOPs — CMA-tuned best config per model family



— NDM (E88) — nonlinear delta, matrix state
— FLA-GDN — linear gated delta-net

B. FLOPs-per-bit rate — four families collapse to one curve



— Mamba2 — linear selective SSM
— E1 — vanilla nonlinear Elman, dense W_h